

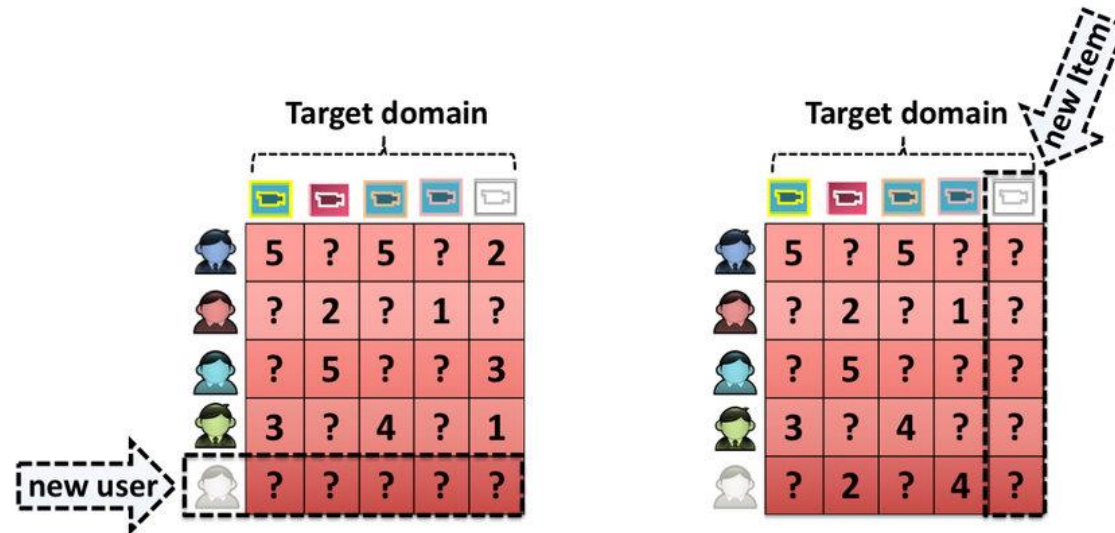
A Review of Cold Start Problem in Recommendation System



Cold Start Problem

- Cold-start (CS) problem: a problem to generate recommendations for a fresh user or to recommend a new item.
- New users, new items, new community (not our focus).
- CS problem, pure CS problem.

pure CS problem: a sub-task of the CS problem, where users have just started to iterate with the systems without any personal information associated.




Categories of Solutions

- Integrate side information:
 - 1. demographic features; 2. product descriptions; 3. cross-domain recommendation.
- Group recommendation: new users/items can be related to the group they belong to
- Examples: PreHash, AGRRE, HIM
- Meta-learning: learning-to-learn process.
- Examples: MeLU, MAMO



Cross-Domain Recommendation

- SSCDR: Semi-Supervised Learning for Cross-Domain Recommendation to Cold-Start Users, 2019.
- Find a mapping $f_{\theta}(u_i^s) = u_i^t$, where u and v are the embedding vectors for users and items, respectively, and OU is the overlapping users set.



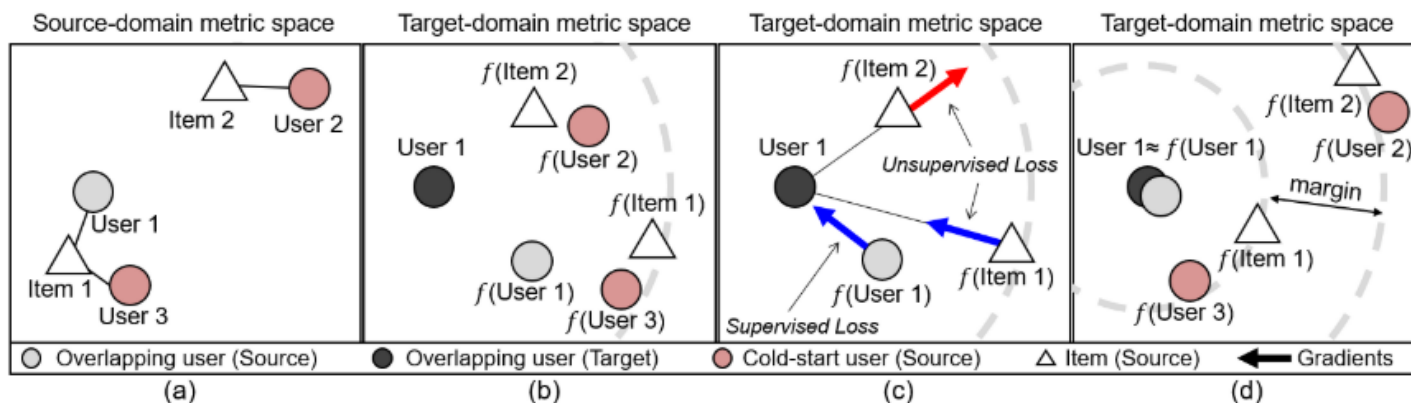
$$\min_{\theta} \mathcal{L}_S + \lambda \cdot \mathcal{L}_U$$

$$s.t. \quad ||f_{\theta}(\mathbf{u}_*)||^2 \leq 1 \text{ and } ||f_{\theta}(\mathbf{v}_*)||^2 \leq 1.$$

- Supervised loss + Unsupervised loss (negative samples):

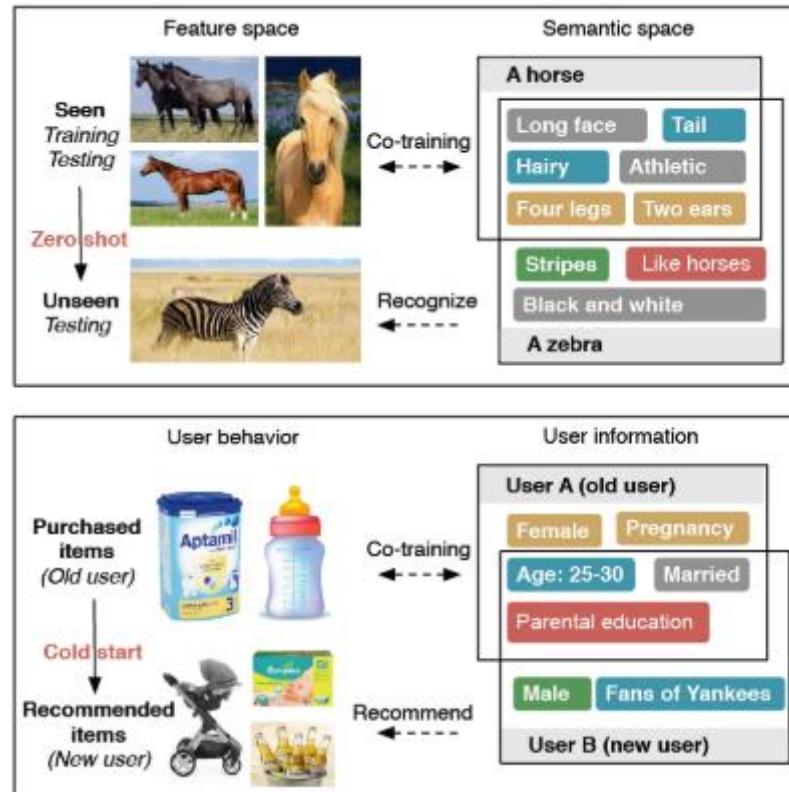
$$\mathcal{L}_S = \sum_{i \in OU} d(f_{\theta}(\mathbf{u}_i^s), \mathbf{u}_i^t),$$

$$\mathcal{L}_U = \sum_{i \in OU} \sum_{j \in NI_i^s} \sum_{k \notin NI_i^s} \left[m + d(f_{\theta}(\mathbf{v}_j^s), \mathbf{u}_i^t) - d(f_{\theta}(\mathbf{v}_k^s), \mathbf{u}_i^t) \right]_+,$$



Zero-Shot Learning

- From Zero-Shot Learning to Cold-Start Recommendation, 2019
- Zero-shot learning (top) and cold-start (bottom)

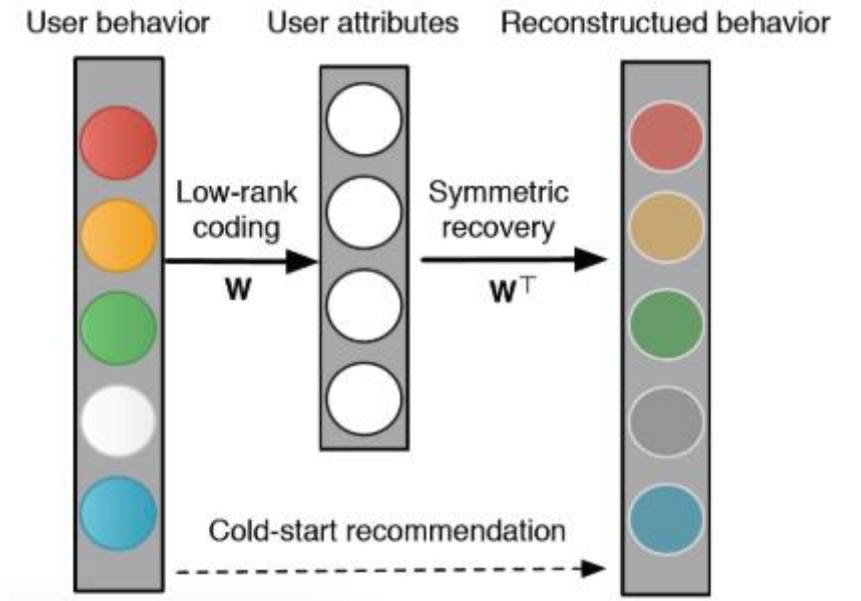


Zero-Shot Learning

- From Zero-Shot Learning to Cold-Start Recommendation, 2019
- X : user behavior
- S : user attributes, like personal information and social network data

$$\min_{\mathbf{W}} \|\mathbf{X} - \mathbf{W}^T \mathbf{W} \mathbf{X}\|_F^2 + \beta \text{rank}(\mathbf{W}), \quad s.t. \mathbf{W} \mathbf{X} = \mathbf{S},$$

- New user: $X_{new} = W^T S_{new}$
- Here a linear auto-encoder is used. Maybe neural auto-encoder will perform better.



Multi-Feature Fusion

- MFDCF: Multi-Feature Discrete Collaborative Filtering for Fast Cold-Start Recommendation, 2020.
- Multi-Feature Fusion:
- Denote user's multiple content features as $X^{(m)}|_{m=1}^M$, where $X^{(m)} = [x_1^{(m)}, \dots, x_n^{(m)}] \in R^{d_m \times n}$.
- Map $X^{(m)}|_{m=1}^M$ into a multi-feature representation $H \in R^{r \times n}$ with low-rank constraint on W .

$$\min_{\mu \in \Delta_M, \mathbf{W}^{(m)}, \mathbf{H}} \sum_{m=1}^M \frac{1}{\mu^{(m)}} \|\mathbf{H} - \mathbf{W}^{(m)} \mathbf{X}^{(m)}\|_F^2 + \gamma \text{rank}(\mathbf{W}^{(m)})$$



Multi-Feature Fusion

- Collaborative Filtering: get hash codes (embedding) for user and item
- user-item rating matrix $S \in R^{n \times m}$; $b_i \in \{\pm 1\}^r$ denote the binary hash codes for i-th user; $d_j \in \{\pm 1\}^r$ denote the binary hash codes for j-th item;
- $B = [b_1, \dots, b_n]$; $D = [d_1, \dots, d_m]$

$$\min_{B,D} \|S - B^T D\|_F^2$$
$$s.t. B \in \{\pm 1\}^{r \times n}, D \in \{\pm 1\}^{r \times m}$$



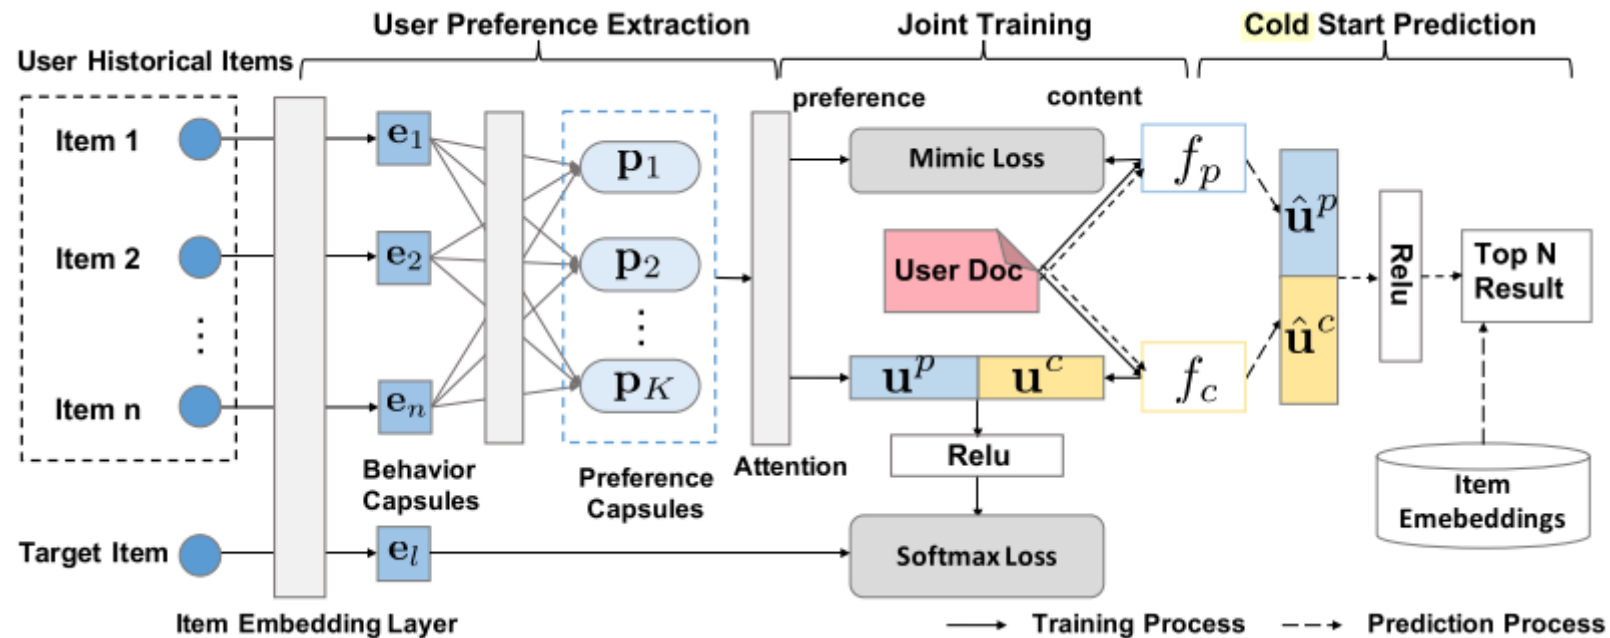
$$\min_{B,D,R} \|S - H^T R^T D\|_F^2 + \beta \|B - RH\|_F^2$$
$$s.t. B \in \{\pm 1\}^{r \times n}, D \in \{\pm 1\}^{r \times m}, R^T R = I_r$$

- For new user, we use $X^{(m)}$ to infer B .
- Linear function can be replaced by more complex function like neural network.
- $W^{(m)} X^{(m)} \rightarrow f(X^{(m)}); B^T D \rightarrow f(B, D)$
- The fusion process can be replaced by mixture-of-experts model.



Capsule Network

- Joint Training Capsule Network for Cold Start Recommendation, 2020.



Group recommendation

- Preference Hash (PreHash): Beyond User Embedding Matrix: Learning to Hash for Modeling Large-Scale Users in Recommendation, 2020.

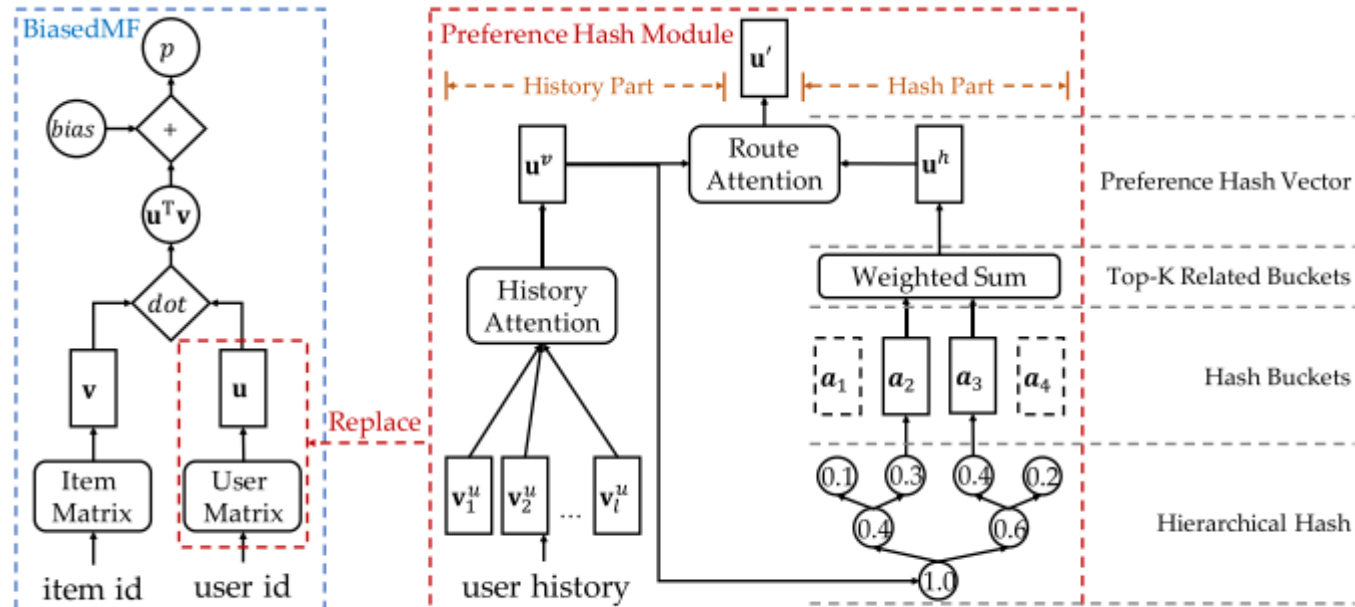


Figure 1: An example of BiasedMF_{PreHash} (PreHash enhanced BiasedMF).



Group recommendation

- In PreHash, the hash vector for new user is a weighted sum of top K related buckets.

$$\mathbf{u}^h = \frac{1}{\sum_{j=1}^K r_{t_j}} \sum_{j=1}^K r_{t_j} \mathbf{a}_{t_j}$$

- Potential cons: the selection of K; if K is small, \mathbf{u}^h may be sensitive to whether \mathbf{a}_{t_j} is among the top K or not.

- Prob PreHash:

- 1. A automatic clustering during the process (mixture Gaussian process) →
for new user: $\mathbf{u}^h \sim \sum_{j=1}^K w_j N(\mu_j, \Sigma_j)$
- 2. User history can be divided into several sessions like long and short term history.
- 3. We can use capsule network to better the embeddings.



Group recommendation

- AGRRE: Attentive Group Recommendation, 2018.

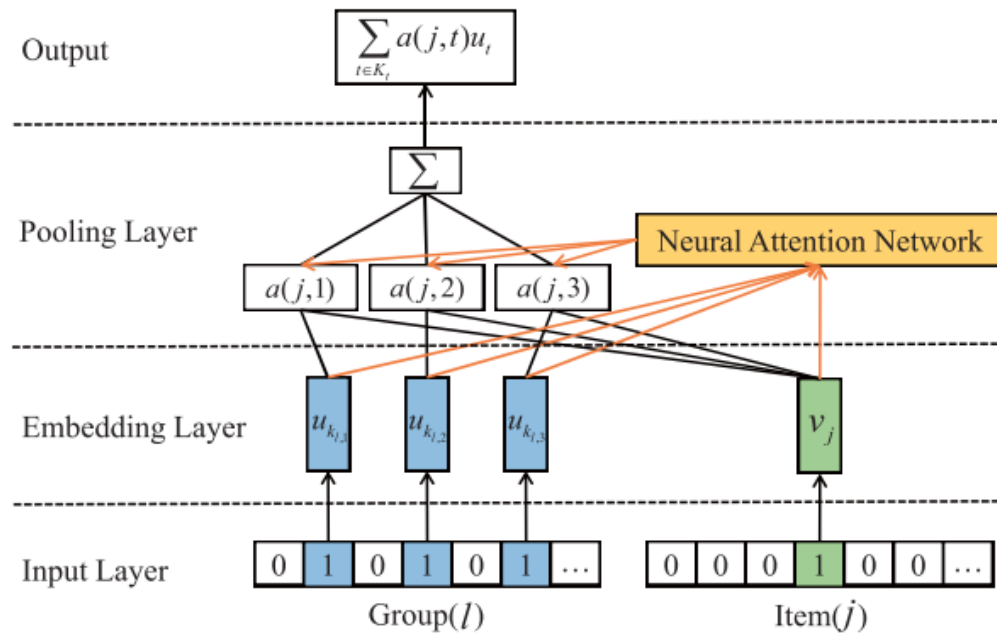


Figure 2: Illustration of the user embedding aggregation component based on neural attention network.



Group recommendation

- AGRRE: Attentive Group Recommendation , 2018.

$$g_l(j) = \underbrace{\sum_{t \in \mathcal{K}_l} \alpha(j, t) \mathbf{u}_t}_{\text{user embedding aggregation}} + \underbrace{\mathbf{q}_l}_{\text{group preference embedding}}$$

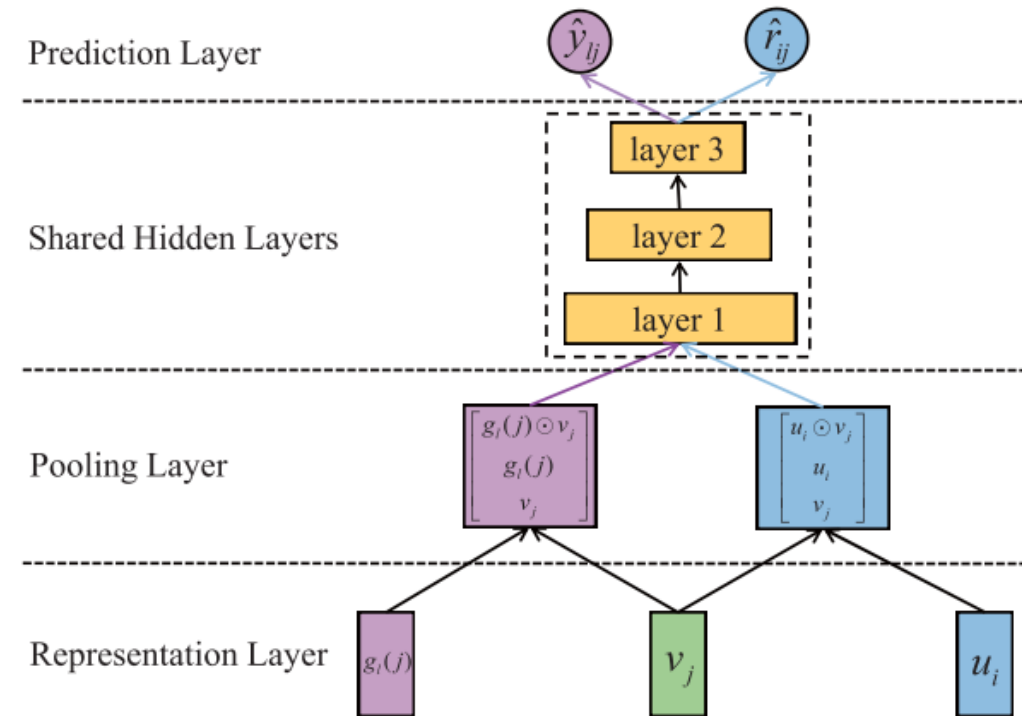


Figure 3: Illustration of interaction learning based on NCF.

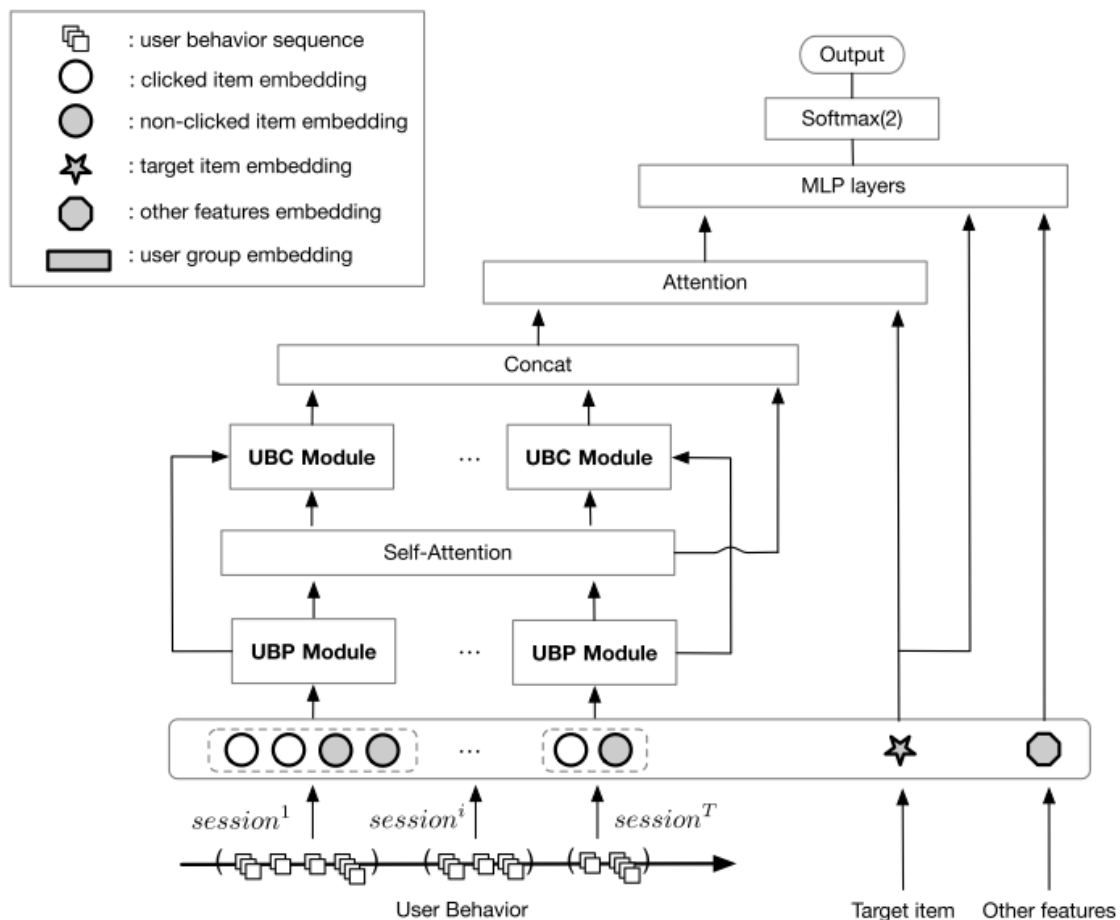


Group recommendation

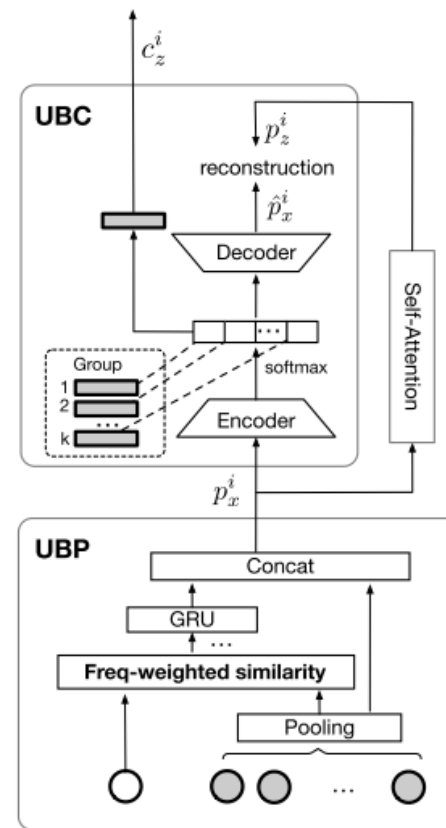
- HIM: Hybrid Interest Modeling for Long-tailed Users, 2020.
- User Behavior Pyramid (UBP) module: to capture the fine-grained personalized interest.
- User Behavior Clustering (UBC) module: to learn latent user interest groups and capture coarse-grained semi-personalized interest.
- Group initialization information is not needed here.



Group recommendation (HIM)



(a) Framework of HIM



(b) Modeling in each session



Meta-learning (learn to learn)

- Traditionally, given $D = \{(x_i, y_i)\}$, we learn $\hat{y} = f_{\theta}(x)$. This usually works when we have lots of data.
- Can we learn new concepts and skills fast with a few training examples?
- Meta-learning: trained over a variety of learning tasks and optimized for the best performance on a distribution of tasks.
- L is a task sample, S^L is a support set, B^L is a prediction set.
- Maximize:

$$\mathbb{E}_{L \sim \mathcal{L}} [\mathbb{E}_{S^L \subset D, B^L \subset D} [\sum_{(\mathbf{x}, y) \in B^L} P_{g_{\phi}(\theta, S^L)}(y|\mathbf{x})]]$$

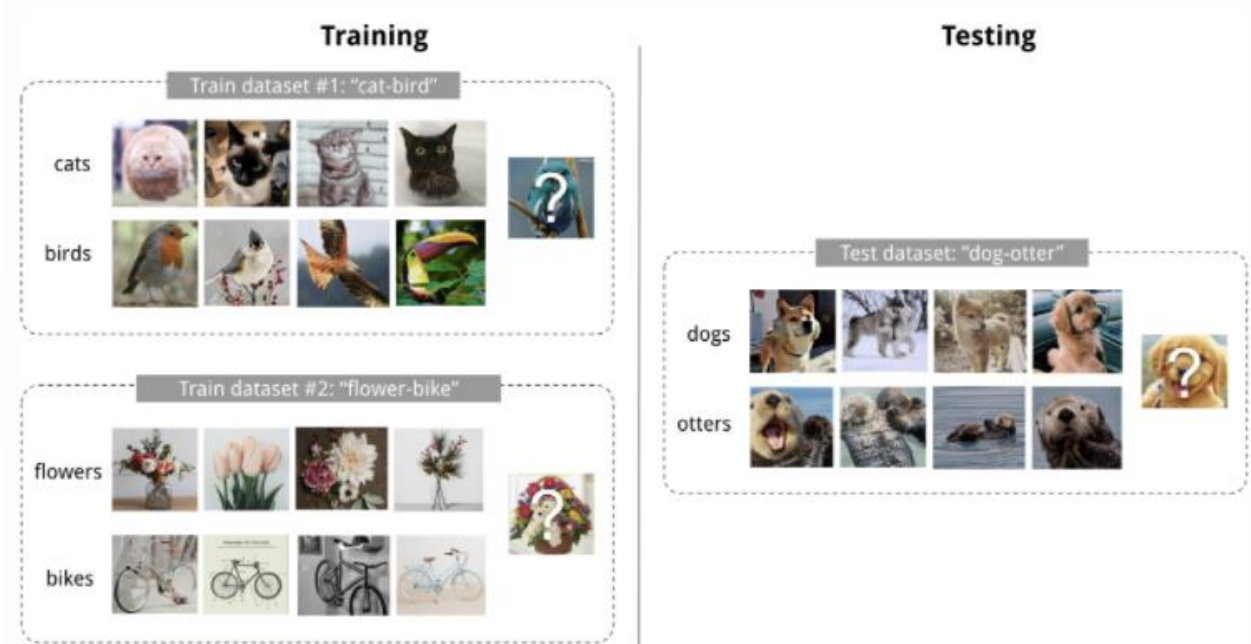



Fig. 1. An example of 4-shot 2-class image classification. (Image thumbnails are from [Pinterest](#))



Meta-learning (optimization-based)

- Optimize the model parameters explicitly for fast learning (good initialization)
- MAML, short for Model-Agnostic Meta-Learning, is compatible with model that learns through gradient descent.
- Local update (ϕ^0 to $\hat{\theta}^m$): sample several tasks.
- Global update (ϕ^0 to ϕ^1): based on query set.



Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks


Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples
 - 6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
 - 7: **end for** Note: the meta-update is using different set of data.
 - 8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
 - 9: **end while**
-



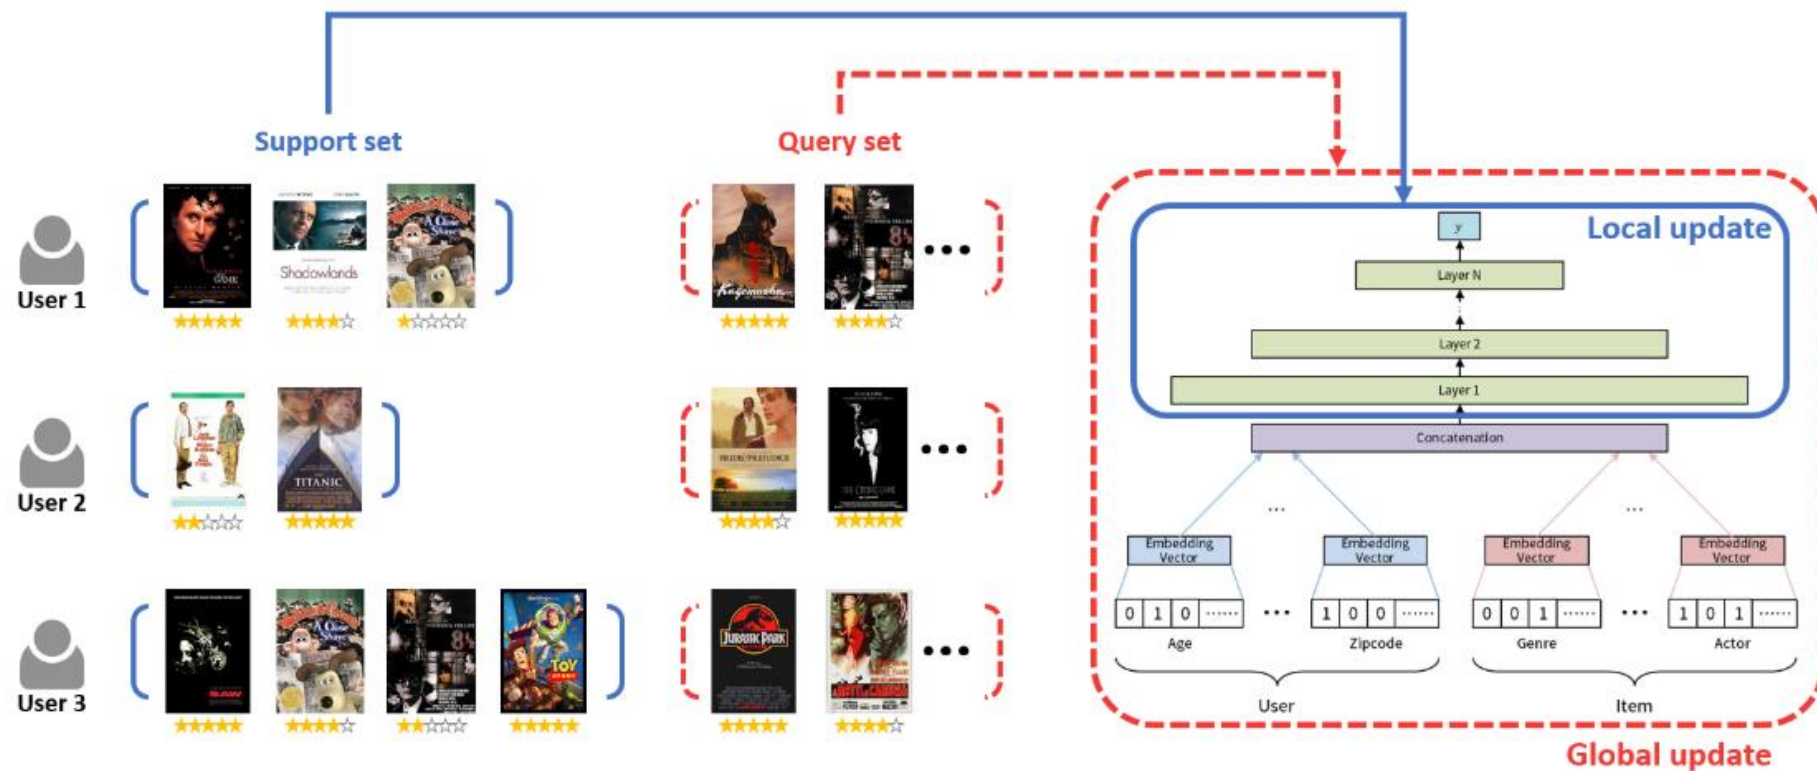
Meta-learning

- MeLU: Meta-Learned User Preference Estimator for Cold-Start Recommendation, 2019.
- MeLU can estimate new user's preferences with a few consumed items (good initialization for the parameters).
- Each user can be regarded as a task.
- Items with large gradient are useful for identifying user preferences.
- We recommend not only products that users like but also products that help us understand users. (balance between exploitation and exploration)
→ Bayesian optimization



Meta-learning

- MeLU: Meta-Learned User Preference Estimator for Cold-Start Recommendation



Meta-learning

- MeLU: Meta-Learned User Preference Estimator for Cold-Start Recommendation

Algorithm 1 Model-Agnostic Meta-Learning for User Preference Estimator


Require: α, β : step size hyperparameters

```
1: randomly initialize  $\theta_1$  (parameters in Eqs. 1 and 2)
2: randomly initialize  $\theta_2$  (parameters in Eq. 3)
3: while not converge do
4:   sample batch of users  $B \sim p(\mathcal{B})$ 
5:   for user  $i$  in  $B$  do
6:     set  $\theta_2^i = \theta_2$ 
7:     evaluate  $\nabla_{\theta_2^i} \mathcal{L}_i(f_{\theta_1, \theta_2^i})$ 
8:     local update  $\theta_2^i \leftarrow \theta_2^i - \alpha \nabla_{\theta_2^i} \mathcal{L}_i'(f_{\theta_1, \theta_2^i})$ 
9:   end for
10:  global update  $\theta_1 \leftarrow \theta_1 - \beta \sum_{i \in B} \nabla_{\theta_1} \mathcal{L}_i'(f_{\theta_1, \theta_2^i})$ 
    $\theta_2 \leftarrow \theta_2 - \beta \sum_{i \in B} \nabla_{\theta_2} \mathcal{L}_i'(f_{\theta_1, \theta_2^i})$ 
11: end while
```



Meta-learning

- MAMO: Memory-Augmented Meta-Optimization for Cold-start Recommendation, 2020
- Global sharing parameter may lead the model into local optima for some users



- 1. Feature-specific memory matrix: guide the model with personalized parameter initialization.
- 2. Task-specific memory matrix: guide the model fast predicting the user preference.



Meta-learning

- MAMO: Memory-Augmented Meta-Optimization for Cold-start Recommendation
- 1. user profile memory M_P + user embedding memory M_U : providing personalized bias term
 - $(P_U, M_U, M_P) \rightarrow b_u$
 - $\theta_u \leftarrow \phi_u - \tau b_u$
 - b_u : guiding the global parameter to fast adapt to the case of user u
- 2. task-specific memory $M_{U,I}$
- We may learn PreHash in a meta-learning scheme.

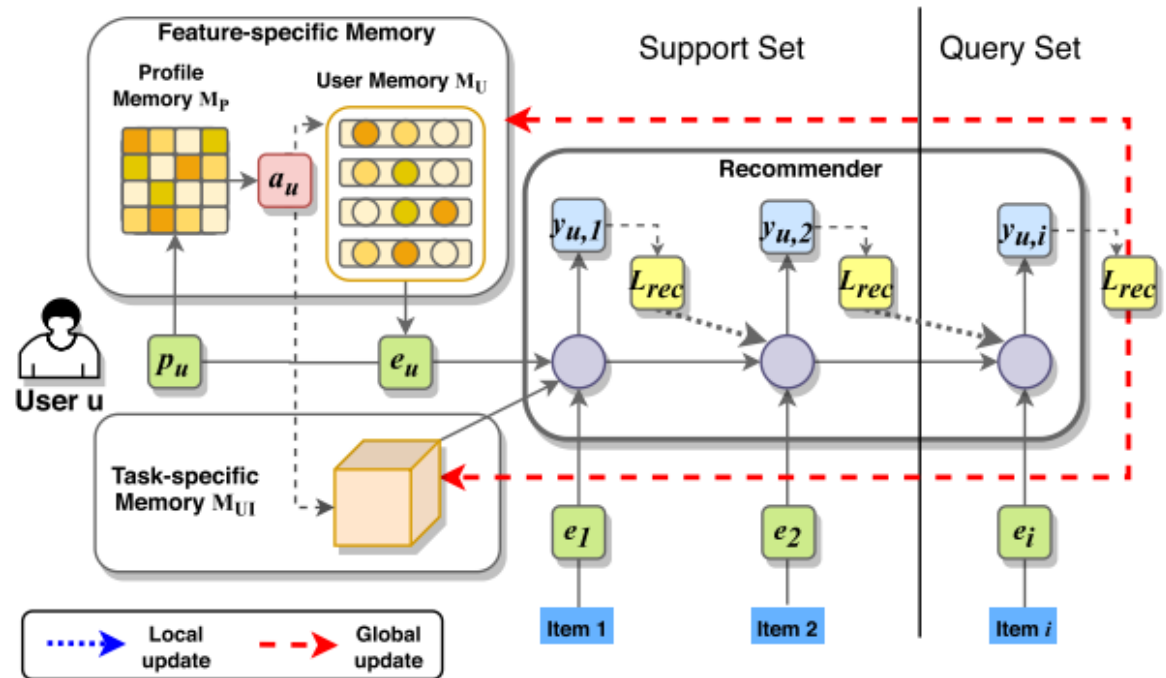


Figure 1: The training phase of MAMO



Pure CS problem

- Pure CS problem: users without historical data
- Knowledge RSs: collect user information using small questionnaires in the first user interaction.
- Social Filtering RSs: external information about users, such as social, demographic and/or personal data.
- Non-personalized RSs: exploit global information about items and users to provide recommendations
- (popularity, recency and positive ratings → max-coverage and category-exploration)

